



Mapping Poverty with Neural Networks

Jean, Burke, Xie, Davis, Lobell, & Ermon (2016)

Graham Joncas (戈雷)

17210680479@fudan.edu.cn

<https://github.com/gjoncas>

School of Economics
Fudan University

Topics in Development Economics
April 4, 2018



Outline

Artificial Neural Networks

Convolutional Neural Nets

Jean et al. (2016) – Background

Methodology & Results


Conclusions



Introduction to Neural Nets

- One of the most commonly used machine learning algorithms
- Many applications in economics – used since White (1988)
- Avoid unrealistic assumptions about data-generating process
- Make non-linear statistics accessible to non-statisticians
- ‘Deep learning’ is just neural nets with lots of hidden layers

Neural Networks & The Brain



- Dendrites: accept inputs
- Axons: transmit output
- If combined inputs meet a threshold, sends an output

“The cerebral cortex contains about 10^{11} neurons, which is approximately the number of stars in the milky way. [...]

Each neuron is connected to 10^3 to 10^4 other neurons. In total, the human brain contains approximately 10^{14} to 10^{15} interconnections.”

— Jain, Mao, Mohiuddin (1996: 33)

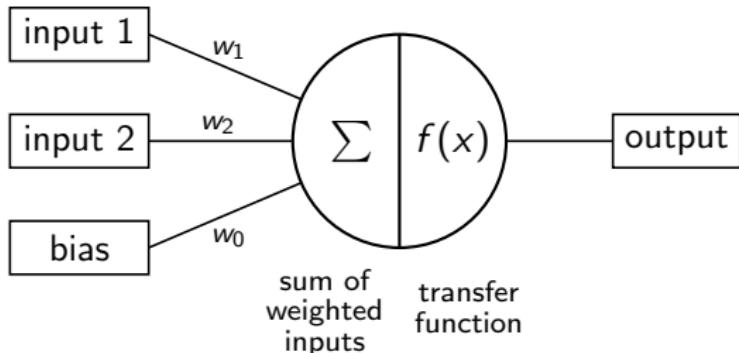


Brain vs. Computer

| | Brain | Computer |
|----------------------------|-------------------------|-------------------------|
| Number of Processing Units | $\approx 10^{11}$ | $\approx 10^9$ |
| Type of Processing Units | Neurons | Transistors |
| Form of Calculation | Massively Parallel | Generally Serial |
| Data Storage | Associative | Address-based |
| Response Time | $\approx 10^{-3}$ s | $\approx 10^{-9}$ s |
| Processing Speed | Very Variable | Fixed |
| Potential Processing Speed | $\approx 10^{13}$ FLOPS | $\approx 10^{18}$ FLOPS |
| Real Processing Speed | $\approx 10^{12}$ FLOPS | $\approx 10^{10}$ FLOPS |
| Resilience | Very High | Almost None |
| Power Consumption per Day | 20W | 300W |



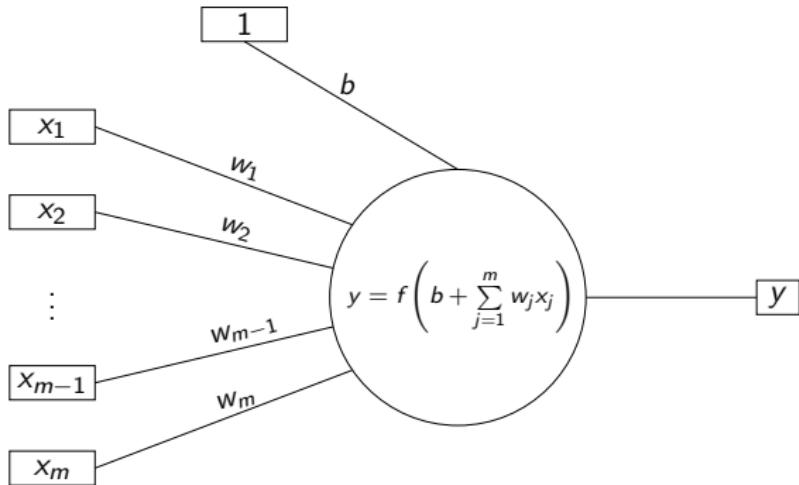
Abstract Model



- Simplest is *perceptron*: binary input and output, one neuron
- Inputs (x_i) have weights (w_i); bias is like β_0 (where $x_0 = 1$)
- Neuron sums weighted inputs ($s = \sum w_i x_i$)
- Transforms sum into output according to threshold θ :

$$f(s) = \begin{cases} 1, & \text{if } s > \theta \\ 0, & \text{otherwise} \end{cases}$$

Machine Learning



- Supervised learning: given correct outputs, learn w_i for inputs
- *Training set* → find weights; *test set* → see if results still hold
- Supervised → pattern classification; unsupervised → clustering

Linear Separability: OR vs. XOR

With linear regression, we want to draw a line that best fits all data points (minimizes $\sum u_i^2$).


Neural nets: want to draw a line that separates different classes

Classes X_0 and X_1 are linearly separable if $\exists \vec{w}, k$ such that:

- For all $x \in X_0$, $\sum_{i=1}^n w_i x_i > k$
- For all $x \in X_1$, $\sum_{i=1}^n w_i x_i < k$

| \vee | 1 | 0 |
|--------|---|---|
| 1 | 1 | 1 |
| 0 | 1 | 0 |

| \vee | 1 | 0 |
|--------|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |





Learning with Neural Nets

Widrow-Hoff Learning Rule

$$\vec{w}_{i+1} = \vec{w}_i + \eta(z_i - y_i)\vec{x}_i$$

Learning the classification for logical OR (\vee)

Use input vectors as training set, including $x_0 = 1$:

$$\vec{x}_1 = \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix}, \vec{x}_2 = \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix}, \vec{x}_3 = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}, \vec{x}_4 = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$$

Also given correct classification: $z_1 = z_2 = z_3 = 1, z_4 = 0$

Perceptron classifies object as 1 iff $w_0x_0 + w_1x_1 + w_2x_2 > 0$, else 0

Use starting vector $\vec{w}_1 = (0, 0, 0)'$, set learning rate $\eta = 1$



Learning Logical OR (\vee)

① \vec{w}_1 gives $y_1 = 0 \neq z_1$. So let's change the weights:

$$\vec{w}_2 = \vec{w}_1 + (z_1 - y_1)\vec{x}_1$$

$$\vec{w}_2 = (0, 0, 0)' + (1 - 0)(1, 1, 0)' = (\textcolor{blue}{1}, \textcolor{blue}{1}, 0)'$$

② \vec{x}_2 is correctly classified with the weight vector \vec{w}_2 .

$$\textcolor{blue}{1} \times 1 + \textcolor{blue}{1} \times 0 + \textcolor{blue}{0} \times 1 = 1 > 0 \rightarrow y_2 = z_2 = 1$$

③ \vec{x}_3 is correctly classified with the weight vector \vec{w}_2 .

$$\textcolor{blue}{1} \times 1 + \textcolor{blue}{1} \times 1 + \textcolor{blue}{0} \times 1 = 2 > 0 \rightarrow y_3 = z_3 = 1$$



Learning Logical OR (\vee)

- ④ For \vec{x}_4 is $\vec{w}'_2 \vec{x}_4 = 1$, so we have to change the weights again:

$$\vec{w}_3 = \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} + (0 - 1) \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$$

- ⑤ \vec{x}_1 is now used as input and is correctly classified.
⑥ Since $\vec{w}'_3 \vec{x}_2 = 0 \neq z_2$:

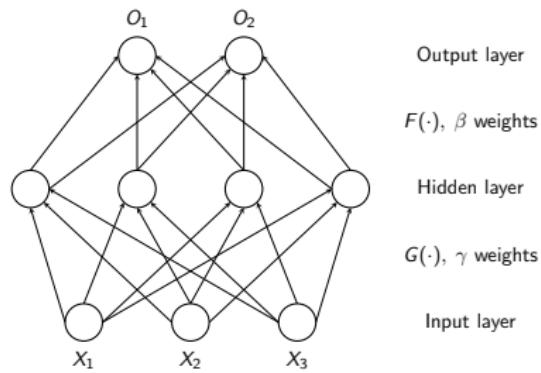
$$\vec{w}_4 = (0, 1, 0)' + (1 - 0)(1, 0, 1)' = (1, 1, 1)'$$

- ⑦ Since $\vec{w}'_4 \vec{x}_3 = 3 > 0$, \vec{x}_3 is correctly classified.
⑧ \vec{x}_4 is incorrectly classified as 1, so that


$$\vec{w}_5 = (1, 1, 1)' + (0 - 1)(1, 0, 0)' = (0, 1, 1)'$$

Our perceptron has learned the correct weights for OR: $(0, 1, 1)'$

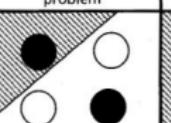
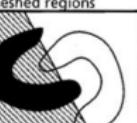
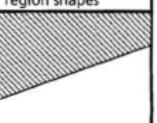



Neural Nets – Feed-forward vs. Recurrent Learning



- ‘Hidden layers’ are nodes not connected to input or output
- Feed-forward: only goes forward
- Recurrent: feedback into inputs




Hyperplane Separation

| Structure | Description of decision regions | Exclusive-OR problem | Classes with meshed regions | General region shapes |
|--------------|--|---|---|--|
| Single layer | Half plane bounded by hyperplane |  |  |  |
| Two layer | Arbitrary (complexity limited by number of hidden units) |  |  |  |
| Three layer | Arbitrary (complexity limited by number of hidden units) |  |  |  |

- Adding hidden layers lets us draw more complex shapes
- General case: search for hyperplane to separate each class
- But: adding too many hidden layers just memorizes the data




Overfitting



More complexity does not mean better prediction

- Blue function is 10th-degree polynomial with exact fit
- Black function is OLS regression line
- New data likely closer to regression line than fitted line

Dealing with Overfitting



- Training set: infer weights
- Validation set: check if \vec{w} holds outside training data

Trade-off between bias & variance:

- model bias - systematic error in learning underlying relations among variables
- model variance - how well a model generalizes elsewhere

Regression may misrepresent true functional form (high bias), but generalize well (low variance)

Neural net may depend too much on data (low bias), generalize poorly (high variance)



Analogies with Regression

Linear Regression

$$y = \beta_0 + \beta_1 x_1 + \dots + \beta_n x_n + \varepsilon$$

Neural Net

$$y = f \left(b + \sum_{i=1}^n w_i \cdot x_i \right)$$

Neural Net model

Single linear perceptron

Simple linear perceptron

Simple nonlinear perceptron

Adaline network

Functional link network

Multilayer perceptron

Statistical equivalent

Linear regression

Multivariate multiple linear regression

Logistic regression


Linear discriminant function

Polynomial regression

Simple nonlinear or multivariate multiple nonlinear regression



Activation Functions





Backpropagation Algorithm

Based on Widrow-Hoff learning rule, extended for multiple neurons

$$w_{ij}^{new} = w_{ij}^{old} + \Delta w_{ij}, \text{ where } \Delta w_{ij} = -\eta \frac{\partial E}{\partial w_{ij}}$$

- ① Start with random weights, forward-propagate to get output
- ② Get error as difference from true answer: $E_i = (F_i(x, w) - z_i)^2$
- ③ Get partial derivatives of error w.r.t. each w_{ij} , add up for total
 - Use sigmoid function because derivative is easy
- ④ Update weights by equation above, repeat until E minimized



Convolutional Neural Nets – Computer Vision



What We See

| | | | | | | | | | | | | | | | | | | | | |
|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 08 | 02 | 22 | 97 | 38 | 15 | 00 | 40 | 00 | 75 | 04 | 05 | 07 | 78 | 52 | 12 | 50 | 77 | 91 | 08 | |
| 49 | 49 | 99 | 40 | 17 | 81 | 18 | 57 | 60 | 87 | 17 | 40 | 98 | 43 | 69 | 48 | 04 | 56 | 62 | 00 | |
| 81 | 49 | 31 | 73 | 55 | 79 | 14 | 29 | 93 | 71 | 40 | 47 | 53 | 88 | 30 | 03 | 49 | 13 | 36 | 65 | |
| 52 | 70 | 95 | 23 | 04 | 60 | 11 | 42 | 69 | 24 | 48 | 56 | 01 | 32 | 54 | 71 | 37 | 02 | 36 | 91 | |
| 22 | 31 | 16 | 73 | 51 | 67 | 69 | 89 | 41 | 92 | 36 | 54 | 22 | 40 | 40 | 28 | 66 | 33 | 13 | 80 | |
| 24 | 47 | 32 | 60 | 99 | 03 | 45 | 09 | 44 | 75 | 33 | 53 | 79 | 36 | 84 | 20 | 35 | 17 | 12 | 50 | |
| 32 | 98 | 81 | 28 | 64 | 23 | 67 | 10 | 26 | 38 | 40 | 67 | 59 | 54 | 70 | 66 | 18 | 38 | 64 | 70 | |
| 67 | 26 | 20 | 68 | 02 | 62 | 12 | 20 | 95 | 63 | 94 | 39 | 63 | 08 | 40 | 91 | 66 | 49 | 94 | 23 | |
| 24 | 55 | 58 | 05 | 62 | 73 | 99 | 29 | 97 | 17 | 78 | 78 | 93 | 83 | 14 | 88 | 34 | 89 | 63 | 72 | |
| 21 | 36 | 23 | 09 | 75 | 00 | 76 | 44 | 20 | 45 | 35 | 14 | 00 | 41 | 33 | 97 | 34 | 31 | 33 | 95 | |
| 78 | 17 | 53 | 28 | 22 | 75 | 31 | 67 | 15 | 94 | 03 | 80 | 04 | 42 | 16 | 14 | 09 | 53 | 56 | 92 | |
| 16 | 39 | 05 | 42 | 94 | 35 | 31 | 47 | 55 | 58 | 88 | 24 | 00 | 17 | 54 | 24 | 36 | 29 | 85 | 57 | |
| 86 | 54 | 00 | 48 | 35 | 71 | 89 | 07 | 05 | 44 | 46 | 37 | 44 | 40 | 21 | 51 | 54 | 17 | 58 | | |
| 19 | 80 | 81 | 68 | 05 | 94 | 47 | 69 | 28 | 73 | 92 | 13 | 84 | 32 | 17 | 77 | 04 | 89 | 55 | 40 | |
| 04 | 52 | 08 | 83 | 97 | 35 | 99 | 14 | 07 | 97 | 57 | 32 | 16 | 26 | 26 | 79 | 33 | 27 | 98 | 66 | |
| 88 | 36 | 68 | 87 | 57 | 62 | 20 | 72 | 03 | 46 | 33 | 67 | 44 | 55 | 12 | 32 | 63 | 93 | 53 | 69 | |
| 04 | 42 | 16 | 73 | 38 | 57 | 25 | 39 | 11 | 24 | 94 | 72 | 18 | 08 | 46 | 29 | 32 | 40 | 62 | 76 | 36 |
| 20 | 69 | 36 | 45 | 72 | 30 | 23 | 84 | 34 | 62 | 99 | 69 | 82 | 67 | 59 | 85 | 74 | 04 | 36 | 16 | |
| 20 | 73 | 35 | 29 | 78 | 31 | 90 | 01 | 74 | 31 | 49 | 71 | 48 | 86 | 81 | 16 | 23 | 57 | 05 | 54 | |
| 01 | 70 | 54 | 71 | 83 | 51 | 54 | 69 | 16 | 92 | 33 | 48 | 61 | 43 | 52 | 01 | 89 | 19 | 47 | 48 | |

What Computers See

- Need numerical inputs → convert image to numbers
- Convert into matrix of RGB values (0–255) for each pixel
- Now we can perform matrix operations on images



Convolution

| | | | | |
|-----------------|-----------------|-----------------|---|---|
| 1 $\times 1$ | 1 $\times 0$ | 1 $\times 1$ | 0 | 0 |
| 0 $\times 0$ | 1 $\times 1$ | 1 $\times 0$ | 1 | 0 |
| 0 $\times 1$ | 0 $\times 0$ | 1 $\times 1$ | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |
| | | |

Convolved
Feature

- Green matrix represents matrix of image (each pixel 0 or 1)
- Yellow ‘filter’ multiplies element-wise at each position
- Output matrix is ‘feature map’, can identify edges, curves, etc.



Convolution

| | | | | |
|---|-----------------|-----------------|-----------------|---|
| 1 | 1 $\times 1$ | 1 $\times 0$ | 0 $\times 1$ | 0 |
| 0 | 1 $\times 0$ | 1 $\times 1$ | 1 $\times 0$ | 0 |
| 0 | 0 $\times 1$ | 1 $\times 0$ | 1 $\times 1$ | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|---|--|
| 4 | 3 | |
| | | |
| | | |
| | | |

Convolved
Feature

- Green matrix represents matrix of image (each pixel 0 or 1)
- Yellow ‘filter’ multiplies element-wise at each position
- Output matrix is ‘feature map’, can identify edges, curves, etc.



Convolution

| | | | | |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|---|---|
| 4 | 3 | 4 |
| | | |
| | | |
| | | |

Convolved
Feature

- Green matrix represents matrix of image (each pixel 0 or 1)
- Yellow ‘filter’ multiplies element-wise at each position
- Output matrix is ‘feature map’, can identify edges, curves, etc.



Convolution

| | | | | |
|-----------------|-----------------|-----------------|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 $\times 1$ | 1 $\times 0$ | 1 $\times 1$ | 1 | 0 |
| 0 $\times 0$ | 0 $\times 1$ | 1 $\times 0$ | 1 | 1 |
| 0 $\times 1$ | 0 $\times 0$ | 1 $\times 1$ | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | | |
| | | |
| | | |

Convolved
Feature

- Green matrix represents matrix of image (each pixel 0 or 1)
- Yellow ‘filter’ multiplies element-wise at each position
- Output matrix is ‘feature map’, can identify edges, curves, etc.



Convolution

| | | | | |
|---|-----------------|-----------------|-----------------|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 _{x1} | 1 _{x0} | 1 _{x1} | 0 |
| 0 | 0 _{x0} | 1 _{x1} | 1 _{x0} | 1 |
| 0 | 0 _{x1} | 1 _{x0} | 1 _{x1} | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | |
| | | |
| | | |

Convolved
Feature

- Green matrix represents matrix of image (each pixel 0 or 1)
- Yellow ‘filter’ multiplies element-wise at each position
- Output matrix is ‘feature map’, can identify edges, curves, etc.

Convolution

| | | | | |
|---|---|-----------------|-----------------|-----------------|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 _{x1} | 1 _{x0} | 0 _{x1} |
| 0 | 0 | 1 _{x0} | 1 _{x1} | 1 _{x0} |
| 0 | 0 | 1 _{x1} | 1 _{x0} | 0 _{x1} |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| | | |

Convolved
Feature

- Green matrix represents matrix of image (each pixel 0 or 1)
- Yellow ‘filter’ multiplies element-wise at each position
- Output matrix is ‘feature map’, can identify edges, curves, etc.



Convolution

| | | | | |
|-----------------|-----------------|-----------------|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 $\times 1$ | 0 $\times 0$ | 1 $\times 1$ | 1 | 1 |
| 0 $\times 0$ | 0 $\times 1$ | 1 $\times 0$ | 1 | 0 |
| 0 $\times 1$ | 1 $\times 0$ | 1 $\times 1$ | 0 | 0 |

Image

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | | |

Convolved
Feature

- Green matrix represents matrix of image (each pixel 0 or 1)
- Yellow ‘filter’ multiplies element-wise at each position
- Output matrix is ‘feature map’, can identify edges, curves, etc.



Convolution

| | | | | |
|---|-----------------|-----------------|-----------------|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 _{x1} | 1 _{x0} | 1 _{x1} | 1 |
| 0 | 0 _{x0} | 1 _{x1} | 1 _{x0} | 0 |
| 0 | 1 _{x1} | 1 _{x0} | 0 _{x1} | 0 |

Image

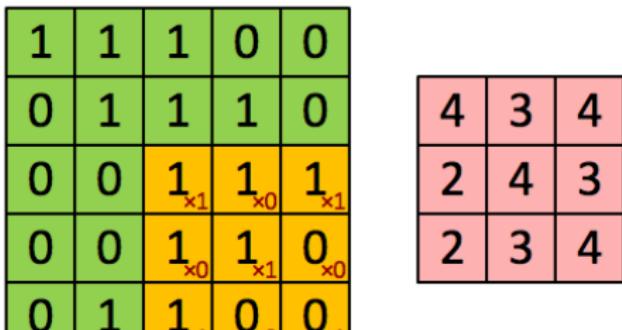
| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | |

Convolved
Feature

- Green matrix represents matrix of image (each pixel 0 or 1)
- Yellow ‘filter’ multiplies element-wise at each position
- Output matrix is ‘feature map’, can identify edges, curves, etc.



Convolution

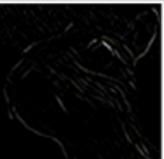
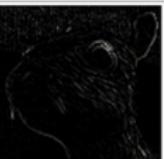


Image

Convolved
Feature

- Green matrix represents matrix of image (each pixel 0 or 1)
- Yellow 'filter' multiplies element-wise at each position
- Output matrix is 'feature map', can identify edges, curves, etc.

Filters & Matrices

| | |
|--|---|
| Identity $\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$ |  |
| Edge detection | $\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$  |
| | $\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$  |
| | $\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$  |

Different filter matrices identify different types of features

- e.g. edges, curves, colors, sharpening, blurring

Convolution operation captures local dependencies in image

- translation invariant: output strength depends not on where features are located, but only that they're present
- e.g. CNN will recognize cats, even if in different positions



Jean et al. (2016) – Motivation

- Very few surveys we can use to construct poverty measures
- World Bank's Living Standards Measurement Study: 39/59 African countries had fewer than two surveys, 14/59 had none
- Demographic & Health Survey: 20/59 African countries had none, 19 more only had one — cost is ~1 million USD each
- Surveys resisted by governments (to hide poor performance)
- Nightlights: hard to distinguish between poor, densely populated areas and wealthy, sparsely populated areas
- Mobile phone data relies on “disparate proprietary datasets”



Jean et al. (2016) – Structure

- ① Pre-train neural net on ImageNet (database of labeled images)
- ② Train neural net to predict nightlight intensities (DMSP data)
based on daytime satellite imagery (Google Maps)

“the model learns to ‘summarize’...high-dimensional input
daytime images as a lower-dimensional set of image features
...predictive of variation in nightlights” (2016: 791)
- ③ Train ridge regressions using cluster-level values from surveys (Y) + image features extracted from daylight imagery (X)
→ helps prevent overfitting on small training set



Pre-training on ImageNet

Recall: Backpropagation Algorithm

$$w_{ij}^{new} = w_{ij}^{old} + \Delta w_{ij}, \text{ where } \Delta w_{ij} = -\eta \frac{\partial E}{\partial w_{ij}}$$

Pre-training as ‘prior distribution’

Before, we used starting weights \vec{w}_1 , which were just random.

- ImageNet: database of labeled images from 1000 categories
- Through classifying images, “the model learns to identify low-level image features such as edges & corners” (2016: 791)
- Further layers learn high-level features, e.g. textures & objects
- Scale, centering differ from satellite images — but good start

Transfer Learning


Daytime satellite photos capture details of the landscape



Convolutional Neural Network (CNN) associates features from daytime photos with nightlight intensity



Satellite nightlights are a proxy for economic activity




Daytime satellite images can be used to predict regional wealth

Household survey locations

CNN processes satellite photos of each survey site

Features from multiple photos are averaged



- CNN model has over 55 million parameters, 8 layers
- Large size due to pixels $400 \times 400 (= 160,000)$, etc.
- Problem: surveys covers a few hundred locations, deep learning requires much more data → transfer learning



Data

- World Bank's LSMS → index of consumption expenditure
- DHS → wealth index, e.g. bikes, TVs, housing material
- Noise added to locations due to privacy ($\pm 5\text{km}$) → clusters
- Countries: Nigeria, Tanzania, Uganda, Malawi, Rwanda
- Nightlights: 400×400 -pixel at zoom level 16 ($\sim 1\text{km}^2$)
- Based on histogram, divide nightlight intensity into 3 classes
- Daylight: for each cluster, up to 100 images; 10km^2 area

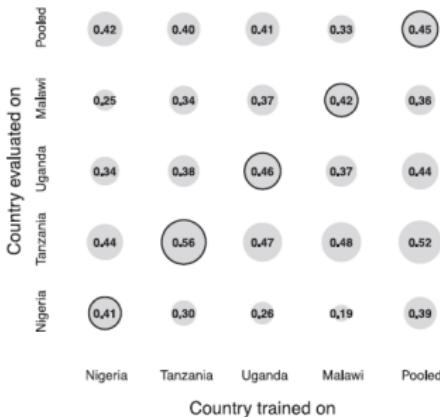


Model learns meaningful features ('filters') on its own




Ridge Regressions

A Consumption expenditures




B Assets



- Pick 100 most useful image features, use as X 's in regression
- Benefit: coefficients are easy to interpret across countries
- Model generalization: β 's same across countries \rightarrow low MSE
- Feature gen: same X 's relevant (β 's significant) \rightarrow high R^2

Transfer Learning vs. Nightlights



- Better at predicting assets (right) than consumption (left)
- R^2 around 37-55% for consumption, 55-75% for asset wealth
- Pooled model (trained across all countries) has R^2 of 44-59%

Extensions: Multi-Resolution Images



- New dataset: Visible Infrared Imaging Radiometer Suite (VIIRS)
- Use images with zoom level 14, 16, 18 → different features
- Low-res images generalize better across countries (same β 's)
- High-res images capture more variation (high R^2 , diff β 's)



Conclusions

"Why should the financial services industry, where mere dollars are at stake, be using more advanced technologies than the aid industry, where human life is at stake?" ~Sendhil Mullainathan



- Machine learning methods far cheaper than official surveys
- Use cutting-edge tools from computer science (deep learning, computer vision) to solve problems in development economics
- Applies out-of-country → use for countries without surveys
- Policy: ensure that resources reach those with greatest need
- Can perhaps be extended for other UN Development Indexes



Machine Learning & 'Patadata'

'Patadata - take metadata output from one model, use as data for other model

- Train neural net on satellite data
- Predict location of borders
- Find exceptions: where neural net says borders *should* be, but aren't



Discrepancy must be due to unobservable factors — institutions, history, culture



References

- Jean, N., Burke, M., Xie, M., Davis, W., Lobell, D., Ermon, S. (2016). “Combining satellite imagery and machine learning to predict poverty.” *Science* 353(6301), pp. 790-4 & S14-32
- Kim, J., Xie, M., Jean, N., Ermon, S. (2016). “Incorporating Spatial Context and Fine-grained Detail from Satellite Imagery to Predict Poverty.” Working Paper. Retrieved from http://cs.stanford.edu/~eix/multi_resolution.pdf
- Xie, M., Jean, N., Burke, M., Lobell, D., Ermon, S. (2015). “Transfer Learning from Deep Features for Remote Sensing and Poverty.” *Proceedings of the 30th AAAI Conference on Artificial Intelligence*, pp. 1-11